

Deep Reinforcement Learning with Continuous Actions


Simon Ramstedt

simonramstedt@gmail.com

Supervisors: Simone Parisi, Gerhard Neumann




TECHNISCHE
UNIVERSITÄT
DARMSTADT






- ▶ Neural Networks work well with **high dimensional** real world data
- ▶ Reinforcement Learning for sequential decision making proved successful for robot learning
- ▶ Deep Reinforcement Learning successes (at least for discrete actions): AlphaGo and Atari games
- ▶ Continuous actions needed for many control tasks

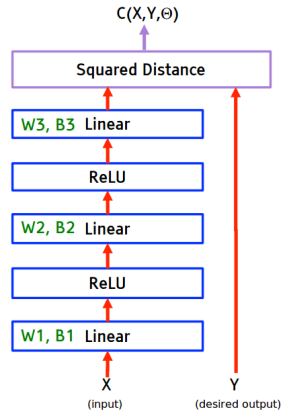
Neural Networks




(a) ReLU: $f(x) = \max(0, x)$.



(b) Tanh: $f(x) = \tanh(x)$.





Objective: Maximize the return (i.e. sum of rewards)

$$R = \sum_{i=t}^T \gamma^{(i-t)} r(s_i, a_i)$$

Policy

$$\pi : s_t \rightarrow a_t$$

Value function

$$V(s_t) = \mathbb{E}[R|s_t]$$

Action value function

$$Q(s_t, a_t) = \mathbb{E}[R|s_t, a_t]$$

Transition model

$$M : s_t, a_t \rightarrow s_{t+1}, r_{t+1}$$


Deep Q-Network

Q-learning with a Neural Network



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Bellman equation: $Q^\pi(s_t, a_t) = \mathbb{E}[r(s_t, a_t) + \gamma Q^\pi(s_{t+1}, \pi(s_{t+1}))]$



- ▶ Problem: correlated samples
- ▶ Solution: experience replay
- ▶ Problem: recursive Q targets
- ▶ Solution: target weights $\theta' = \text{EMA}(\theta)$



for *every timestep* t **do**

 Select action $a_t = \epsilon$ -greedy [$\operatorname{argmax}_a Q(s_t, a|\theta)$]

 Execute a_t and observe reward r_t and next state s_{t+1}

 Store transition (s_t, a_t, r_t, s_{t+1}) in D

 Sample random minibatch of m transitions from D


 Set targets $y_j = r_j + \gamma \max_a Q(a, s_{j+1}|\theta')$

 Perform gradient descent on cost $C = \frac{1}{m} \sum_j (y_j - Q(s_j, a_j|\theta))^2$ with respect to θ

 Update $\theta' \leftarrow \operatorname{LP}(\theta)$

end

Deep Deterministic Policy Gradient (DDPG)



- Policy gradient $\nabla_\xi \mathbb{E}[R] = \nabla_a Q \cdot \nabla_\xi \pi$
- Use tricks from DQN

- 20+ different motor control tasks (pole swingup, cart-pole, double pole, quadruped balance, ...)
- Evaluated on state inputs and pixel inputs
- Smooth reward functions
- Performance: max. **10^6** timesteps of interaction needed



for every timestep t do

Select action $a_t = \pi(s_t) + \mathcal{M}_t$

Execute a_t and observe reward r_t and next state s_{t+1}

Store transition (s_t, a_t, r_t, s_{t+1}) in D

Sample random minibatch of m transitions from D

Set Q targets $y_j = r_j + \gamma Q(s_{j+1}, \pi(s_{j+1}|\xi') | \theta')$

Perform gradient descent on cost $C = \frac{1}{m} \sum_j (y_j - Q(s_j, a_j|\theta))^2$ with respect to θ

Perform gradient ascent on $Q(s_j, \pi(s_j|\xi) | \theta)$ with respect to ξ

Update $\theta' \leftarrow \text{LP}(\theta)$

Update $\xi' \leftarrow \text{LP}(\xi)$

end

Experiments

DDPG sounds nice but ...



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- ▶ How hard is it to implement?
- ▶ How data efficient is it?
- ▶ What about sparse reward functions?
- ▶ How robust is it to hyperparameters variations?
- ▶ Other problems?
- ▶ How can it be improved?


Experiments

Setup




TECHNISCHE
UNIVERSITÄT
DARMSTADT


Policy network layout:



Q network layout:



Cart pole swingup and balance



s_0 : position of the cart

s_1 : angle of the pole


$s_2 = \dot{s}_0$

$s_3 = \dot{s}_1$


a : force applied to the cart

Experiments


Trajectories




(a) Trajectory after 0 timesteps of learning.



(b) Swingup after $\sim 10,000$ timesteps.



(c) Balance after $\sim 100,000$ timesteps.




(d) Near optimal after $\sim 200,000$ timesteps.

Results

DDPG in Cartpole



TECHNISCHE
UNIVERSITÄT
DARMSTADT



$$r = 0.5 \cdot \cos(s_1) - 0.03 \cdot a^2 - 0.015 \cdot |s_0| - 0.2 \cdot s_3^2$$

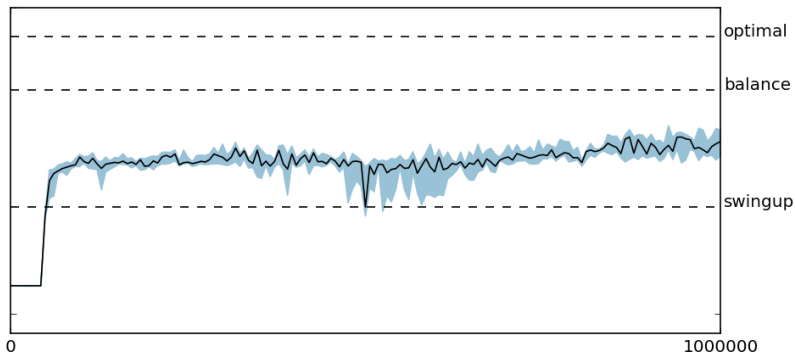
$$t_{\text{warmup}} = 50,000$$

Results

DDPG in Cartpole without target networks



TECHNISCHE
UNIVERSITÄT
DARMSTADT



$$r = 0.5 \cdot \cos(s_1) - 0.03 \cdot \dot{a}^2 - 0.015 \cdot |s_0| - 0.2 \cdot s_3^2$$


$$t_{\text{warmup}} = 50,000$$

Results

DDPG in Cartpole without replay memory



TECHNISCHE
UNIVERSITÄT
DARMSTADT



$$r = 0.5 \cdot \cos(s_1) - 0.03 \cdot \dot{\alpha}^2 - 0.015 \cdot |s_0| - 0.2 \cdot s_3^2$$


$$t_{\text{warmup}} = 50,000$$

Results

DDPG in Cartpole with sparse reward function



TECHNISCHE
UNIVERSITÄT
DARMSTADT



$$r = 1 \cdot (s_1 < 0.01) - 0.03 \cdot a^2$$


$$t_{\text{warmup}} = 10,000$$

Results

DDPG in Cartpole with sparse reward function



TECHNISCHE
UNIVERSITÄT
DARMSTADT




$$r = 4 \cdot (s_1 < 0.01) - 0.03 \cdot a^2$$

$$t_{\text{warmup}} = 10,000$$

Results

DDPG in Cartpole with sparse reward function



$$r = 4 \cdot (s_1 < 0.01) - 0.03 \cdot a^2$$


$$t_{\text{warmup}} = 50,000$$

Results

DDPG in Cartpole with sparse reward function



TECHNISCHE
UNIVERSITÄT
DARMSTADT



$$r = 4 \cdot (s_1 < 0.01) - 0.03 \cdot a^2$$

$$t_{\text{warmup}} = 100,000$$


Future Work

Increase data efficiency



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Share weights:



Learn a **transition model**

Use it to compute another policy gradient (e.g. Heess et. al., 2015)

$$\mathbb{E}[R] = r(s_t, a_t) + \lambda V(M(s_t))$$

$$\nabla_{\xi} \mathbb{E}[R] = \nabla_{\xi} r(s_t, a_t) + \lambda \nabla_{\xi} \pi(s_t) \cdot \nabla_a M(s_t, a_t) \cdot \nabla_{s_{t+1}} V(s_{t+1})$$



Faster reward propagation via **n-step updates** on recent trajectories (e.g. Mnih et. al., 2016)

$$Q^\pi(s_t, a_t) = \mathbb{E}[r_t + \gamma^1 r_{t+1} + \dots + \gamma^n r_{t+n} + \gamma^{n+1} Q^\pi(s_{t+n+1}, a_{t+n+1})]$$

Stochastic neural networks:

- a) Additive noise in the layers (works with standart backpropagation)
- b) Stochastic backpropagation (Kingma & Welling, 2013) (Rezende et. al., 2014)


Thanks

Questions or Feedback ?



TECHNISCHE
UNIVERSITÄT
DARMSTADT

- ▶ Thesis and slides: <https://github.com/simonramstedt/bt>
- ▶ Code: <https://git.ias.informatik.tu-darmstadt.de/SimonR/DRLTF>
- ▶ Deep Reinforcement Learning papers overview:
<https://github.com/junhyukoh/deep-reinforcement-learning-papers>
- ▶ DQN: Mnih et. al., Human-level control through deep reinforcement learning, 2015
- ▶ DDPG: Lillicrap et. al., Continuous control with deep reinforcement learning, 2015
- ▶ A3C: Mnih et. al, Asynchronous Methods for Deep Reinforcement Learning, 2016
- ▶ Prioritized Replay: Schaul et. al., Prioritized Experience Replay, 2015



Deep Q-Network (DQN)

on the Arcade Learning Environment (ALE)

Atari Environment:

- ▶ **Discrete** actions (buttons on Atari controller)
- ▶ Learning from **raw pixel** input
- ▶ Using only the game score as reward signal
- ▶ Performance: max. **10^8** timesteps of interaction needed

Method	Training Time	Mean	Median
DQN (from [Nair et al., 2015])	8 days on GPU	121.9%	47.5%
Gorila [Nair et al., 2015]	4 days, 100 machines	215.2%	71.3%
Double DQN [Van Hasselt et al., 2015]	8 days on GPU	332.9%	110.9%
Dueling Double DQN [Wang et al., 2015]	8 days on GPU	343.8%	117.1%
Prioritized DQN [Schaul et al., 2015]	8 days on GPU	463.6%	127.6%
A3C, FF	1 day on CPU	344.1%	68.2%
A3C, FF	4 days on CPU	496.8%	116.6%
A3C, LSTM	4 days on CPU	623.0%	112.6%

